




Introduction to Data Access and Storage

Thomas Rosenthal – DSI @ UofT

Session 05

Expanding your Database:

INSERT, UPDATE, DELETE

Views

Importing and Exporting Data

Two Final Joins

Expanding your Database:

INSERT, UPDATE, DELETE

Views

Importing and Exporting Data

Two Final Joins

INSERT, UPDATE, DELETE

Prior to this, we've focused solely on retrieving values from tables:

- Tables can also be manipulated with `INSERT`, `UPDATE`, and/or `DELETE`
- *A word of warning...these commands are challenging to undo and can be PERMANENT* 🤖
 - Generally, follow a policy that avoids altering data
 - Make backups of tables before you run a query
 - Never hurts to test on a temporary table first!
- But they are useful, and sometimes the correct solution
- There is no `SELECT` statement for these types of queries

INSERT, UPDATE, DELETE

INSERT

- `INSERT` allows you to add a record
- Specify where you want to add:
 - `INSERT INTO [some_table_name]`
- ...and what you want to add:
 - `VALUES(column_one_value, column_two_value)`
- `VALUES` come in the order of the columns within the tables
- `VALUES` must respect table constraints
 - e.g. NULLs, UNIQUE, data types, etc
- `INSERT` can help create small helper tables
 - **Can we think of any scenarios?**

INSERT, UPDATE, DELETE

UPDATE

- `UPDATE` allows you to change a record
- Specify where you are making your change:
 - `UPDATE [some_table_name]`
- ...and what you want to change:
 - `SET column_one = value1, column_two = value2`
- *SPECIFY A WHERE CONDITION*
 - `WHERE condition`
- You can change a single column, a few columns, all the columns, etc
 - (Respecting table constraints)
- What happens if you don't specify a `WHERE` condition?

INSERT, UPDATE, DELETE

DELETE

- `DELETE` allows you to remove a record
- Specify where you want to delete:
 - `DELETE FROM [some_table_name]`
- *SPECIFY A WHERE CONDITION*
 - `WHERE condition`
- What happens if you don't specify a `WHERE` condition?!?
- `DELETE` doesn't *remove* a table from a database
 - Instead it removes the data from it, leaving the table structure and constraints in place
 - `DROP TABLE` instead if you want to remove it altogether

INSERT, UPDATE, DELETE

(INSERT, UPDATE, DELETE live coding with a TEMP TABLE)

Expanding your Database:

INSERT, UPDATE, DELETE

Views

Importing and Exporting Data

Two Final Joins


Views

- Views instantiate a query result permanently
- They are particularly useful in highly normalized databases, where reproducing a query is tiresome or prone to query errors
- In databases that have live data flowing in:
 - Tables that are created from queries need to be continuously updated whenever there is new data
 - This requires either downtime where the table is empty
 - Or the chance of a "dirty read" (where a table is read before the data is fully updated)
 - Views, on the other hand, will always show the most up-to-date values!
- Views are created just like tables:

```
CREATE VIEW history AS  
SELECT ...
```

Views Performance

- Views can be very slow if poorly created
- Always use primary keys and indexing to make them more performative
- Select the most important columns
- Avoid stacking views (views within a view)
- In commercial RDBMs, use execution plans and/or performance dashboards to analyze the underlying engine mechanisms the view uses for instantiation



Views

(Views live coding)

Expanding your Database:

INSERT, UPDATE, DELETE

Views

Importing and Exporting Data

Two Final Joins

Importing and Exporting Data

Import & Export

CSV

JSON

Importing and Exporting Data

Import & Export

CSV

JSON

Import & Export

- RDBMs allow data to flow into and out of them
 - Some processes are easy:
 - e.g. exporting a table as a CSV file
 - ...while others are complex
 - e.g. writing a CRM to a normalized data warehouse on a nightly basis
- In DB Browser for SQLite, we can make use of the following:
 - Import and export CSV files
 - Manipulate and export JSON files
- SQLite more broadly can:
 - Produce CSVs from queries (using the command line, which we won't do)
 - Connect to other programming languages

Importing and Exporting Data

Import & Export

CSV

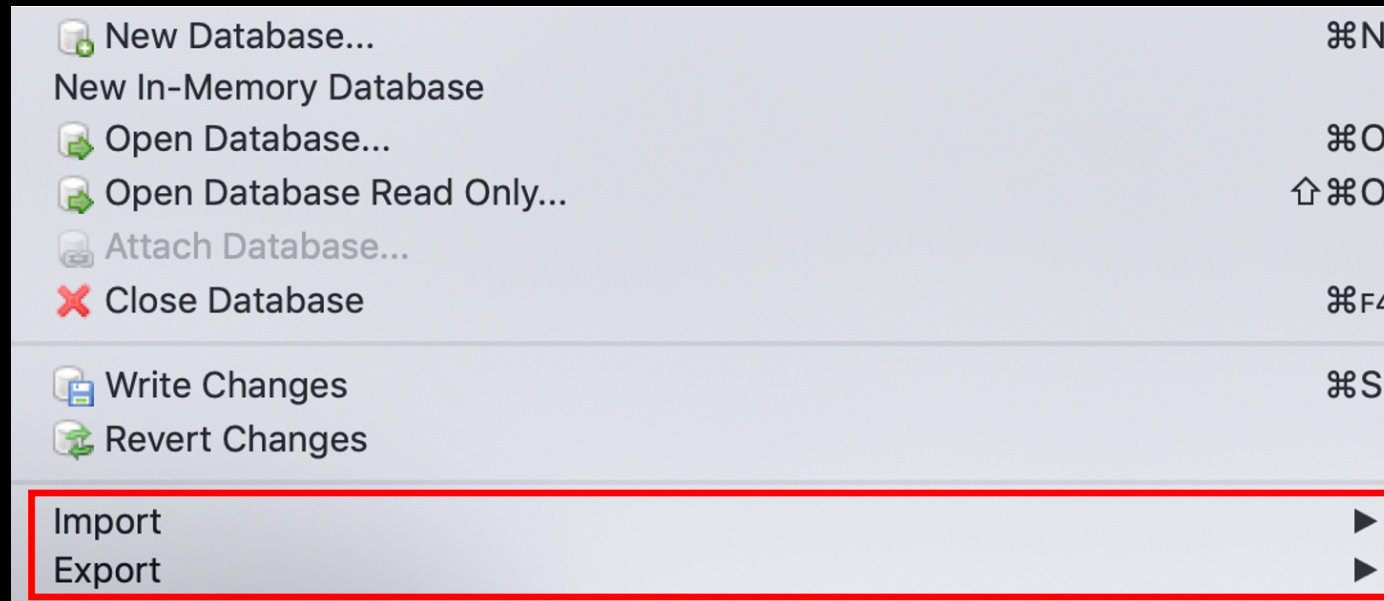
JSON

CSV

- CSV stands for "Comma Separated Values"
- CSVs are file formats well designed to store SQL tables
 - The values of each row are separated by commas
 - Sometimes it makes more sense to use a "|" (pipe), if there is complex text data stored, which might be more sensitive to the presence of commas and/or line breaks
- They are a common file format for transporting structured data
- CSVs can be opened by:
 - Excel
 - Simple text editors (e.g. notepad++, sublime)
 - Programming languages (e.g. python, R)

CSV

DB Browser for SQLite natively supports CSV importing and exporting for tables:



You can also export queries if they are stored in Temporary Tables

CSV

DB Browser for SQLite allows us to extract a query result in a somewhat roundabout method:

- First, write a query

```
SELECT * FROM product p
JOIN product_category pc ON p.product_category_id = pc.product_category_id
```

- Second, right click the results, and select "Copy as SQL"
- Third, instantiate a table with `CREATE`

```
CREATE TABLE "example" ("product_id", "product_name", "product_size",
"product_category_id", "product_qty_type", "product_category_name") ;
```

- Fourth, paste the results from your clipboard

```
INSERT INTO "example" ("product_id", "product_name", "product_size",
"product_category_id", "product_qty_type", "product_category_name")
VALUES ('1', 'Habanero Peppers - Organic', 'medium', '1', 'lbs', 'Fresh Fruits & Vegetables');
...etc for each row
```

- Finally, export the table to CSV

CSV

We can also extract a query result to CSV with python:

```
import pandas as pd
import sqlite3

#set your location, slash direction will change for windows and mac
DB = '/Users/thomas/Documents/GitHub/DSI_SQL/SQL/FarmersMarket.db'

#establish your connection
conn = sqlite3.connect(DB, isolation_level=None,
                        detect_types=sqlite3.PARSE_COLNAMES)

#run your query, use "\"" to allow line breaks
db_df = pd.read_sql_query("SELECT p.*,pc.product_category_name \
                           FROM product p \
                           JOIN product_category pc \
                           ON p.product_category_id = pc.product_category_id"
                           ,conn)

#save
db_df.to_csv('database-py.CSV', index=False)
```


CSV

...or extract with R instead:

```
library(DBI)
library(RSQLite)

#set your location, slash direction will change for windows and mac
DB = '/Users/thomas/Documents/GitHub/DSI_SQL/SQL/FarmersMarket.db'

#establish your connection
conn <- dbConnect(SQLite(), DB)

#run your query
db_df <- dbGetQuery(conn, "SELECT p.*,pc.product_category_name
                           FROM product p
                           JOIN product_category pc
                           ON p.product_category_id = pc.product_category_id")

#save
write.CSV(db_df, file = "database-R.CSV")
```

CSV

(CSV live importing to update our view, CSV live exporting)

Importing and Exporting Data

Import & Export

CSV

JSON

JSON

- JSON stands for "JavaScript Object Notation"
- JSONs are file formats well designed to store tables, lists, arrays, and nested objects
 - Their syntax follows specific rules:
 - Data is in name/value pairs
 - Data is separated by commas
 - Curly brackets '{ }' hold objects
 - Square brackets '[']' hold arrays
 - e.g. `{"first_name":"Ralph", "last_name":"Kimball"}`
 - or for tables:

```
[ {"first_name":"Ralph", "last_name":"Kimball"},  
  {"first_name":"Bill", "last_name":"Imnom"} ]
```

- JSON can be opened by:
 - Web browsers
 - Simple text editors (e.g. notepad++, sublime)
 - Programming languages (e.g. python, R)
- SQLite also provides support for JSON value query and manipulation

JSON

DB Browser for SQLite supports a lot of JSON functions:

- Some are helper functions:
 - `JSON` and `JSON_VALID`, which confirm whether or not a string is JSON and/or in a valid JSON format
 - `JSON_TYPE`
 - When using extracting, type will help to inform column types that SQL will assume, based on the JSON
- Other functions can be used for manipulation or extraction:
- `JSON_EXTRACT` will allow you to return the values of a well-formed JSON string into desired parts
 - Importing JSON into SQL will use either `JSON_EXTRACT` or `JSON_EACH`
 - SQLite gives the following (fairly comprehensive) example set:
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$') → '{ "a":2, "c": [4,5, { "f":7 } ] }'`
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$.c') → '[4,5, { "f":7 } ]'`
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$.c[2]') → '{ "f":7 }'`
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$.c[2].f') → 7`
 - `json_extract('{ "a":2, "c": [4,5], "f":7 }', '$.c', '$.a') → '[[4,5],2]'`
 - `json_extract('{ "a":2, "c": [4,5], "f":7 }', '$.c[#-1]') → 5`
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$.x') → NULL`
 - `json_extract('{ "a":2, "c": [4,5, { "f":7 } ] }', '$.x', '$.a') → '[null,2]'`
 - `json_extract('{ "a": "xyz" }', '$.a') → 'xyz'`
 - `json_extract('{ "a": null }', '$.a') → NULL`

JSON

Importing a JSON array (table) into SQL with DB Browser for SQLite requires a bit more of nuanced approach:

- First copy and paste your JSON table array into SQLite, and put it in a temp table:

```
CREATE TEMP TABLE IF NOT EXISTS temp.[new_json]  
(col BLOB);
```

```
INSERT INTO temp.[new_json](col)  
VALUES('["a": 7, "b": "string"]')
```

- Second, use the `JSON_EACH` function as a table-valued function

```
SELECT key,value  
FROM new_json,JSON_EACH(new_json.col, '$' )
```


- Third, use this previous query as a subquery and combine with `JSON_EXTRACT`, using the value column `JSON_EACH` generated

```
SELECT * ,  
JSON_EXTRACT(value, '$.a') AS a,  
JSON_EXTRACT(value, '$.b') AS b  
  
FROM (...{subquery goes here}...)
```

- You now have a normal SQL table from a JSON array!

JSON

DB Browser for SQLite natively supports JSON exporting for tables:



This also works for Temporary Tables, so queries can be exported as well

JSON

(JSON live exporting)

Expanding your Database:

INSERT, UPDATE, DELETE

Importing and Exporting Data

Views

Two Final Joins

Two Final Joins

CROSS JOIN

Self Joins

CROSS JOIN

- `CROSS JOIN` creates an unfiltered Cartesian Product
- They are not joined on any columns
- Recall our deck of cards example in Session 2:

```
SELECT suit, rank
FROM suits
CROSS JOIN ranks
```

- Because tables 'suits' and 'ranks' contain no common columns, we would have no other means to join
- I love to `CROSS JOIN`!
- They can be super useful when used correctly
 - **What are some good examples that could be useful?** 🧠💬 **Think, Pair, Share**

CROSS JOIN

(CROSS JOIN live coding)

No complex query is complete without at least one CROSS JOIN

(me, jokingly)

Two Final Joins

CROSS JOIN

Self Joins

Self Joins

- Self Joins are somewhat uncommon, but are the last type of possible join
- They are useful for comparison:
 - Determine maximum to-date
 - Generating pairings
- They can help with hierarchy
 - Child-to-Parent relationships
- The syntax is as we might expect:

```
SELECT
e.name as employee_name,
m.name as manager_name

FROM people e
LEFT JOIN people m ON e.manager_id = m.id
```

emp_id	name	manager_id
1	Peter Gibbons	3
2	Michael Bolton	1
3	Bill Lumbergh	
4	Milton Waddams	3

emp_id	name	manager_name
1	Peter Gibbons	Bill Lumbergh
2	Michael Bolton	Peter Gibbons
3	Bill Lumbergh	null
4	Milton Waddams	Bill Lumbergh

Bill is Peter and Milton's boss.

Peter is Michael's boss.    

Self Joins

(Self Joins live coding)

